

Contribution Title

First Author
NTUA

Second Author
NTUA

Third Author
NTUA

Author
NTUA

May 8, 2018

Abstract

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis sapien nunc, commodo et, interdum suscipit, sollicitudin et, dolor. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum. Nam quis nulla. Nullam sit amet magna in magna gravida vehicula. Sed elit dui, pellentesque a, faucibus vel, interdum nec, diam. Mauris dolor felis, sagittis at, luctus sed, aliquam non, tellus. Etiam egestas wisi a erat. Etiam dictum tincidunt diam. Quisque tincidunt scelerisque libero. Phasellus faucibus molestie nisl. Duis sapien nunc, commodo et, interdum suscipit, sollicitudin et, dolor.

Aenean placerat. Duis sapien nunc, commodo et, interdum suscipit, sollicitudin et, dolor. Curabitur ligula sapien, pulvinar a vestibulum quis, facilisis vel sapien. Donec iaculis gravida nulla. Pellentesque arcu. Nulla accumsan, elit sit amet varius semper, nulla mauris mollis quam, tempor suscipit diam nulla vel leo. Nam quis nulla. In convallis. Mauris dolor felis, sagittis at, luctus sed, aliquam non, tellus. Nullam at arcu a est sollicitudin euismod. Sed ac dolor sit amet purus malesuada congue. Integer tempor. Quis autem vel eum iure reprehenderit qui in ea voluptate velit esse quam nihil molestiae consequatur, vel illum qui dolorem eum fugiat quo voluptas nulla pariatur? Vivamus ac leo pretium faucibus. Donec quis nibh at felis congue commodo. Phasellus rhoncus. Duis bibendum, lectus ut viverra rhoncus, dolor nunc faucibus libero, eget facilisis enim ipsum id lacus. Ut enim ad minima veniam, quis nostrum exercitationem ullam corporis suscipit laboriosam, nisi ut aliquid ex ea commodi consequatur?

1 Introduction

1.1 Friedkin-Johnsen Model and Opinion Formation Games

In the Friedkin-Johnsen model (FJ-model) an undirected graph $G(V, E)$ with n nodes, is assumed, where V denotes the agents and E the social relations between them. Each agent i poses an internal opinion $s_i \in [0, 1]$ and a self confidence coefficient $\alpha_i \in [0, 1]$. At each round $t \geq 1$, agent i updates her opinion as follows:

$$x_i(t) = (1 - \alpha_i) \frac{\sum_{j \in N_i} x_j(t-1)}{|N_i|} + \alpha_i s_i, \quad (1)$$

where N_i is the set of her neighbors. The simplicity of the update rule (1) makes the FJ-model plausible, because in real social networks it is very unlikely that agents change their opinions according to complex rules. Based on the FJ-model, in [BKO11], they propose an *opinion formation game* in which the strategy that each agent i plays, is the opinion $x_i \in [0, 1]$ that she publicly expresses incurring her a cost

$$C_i(x_i, x_{-i}) = (1 - a_i) \sum_{j \in N_i} (x_i - x_j)^2 + a_i (x_i - s_i)^2. \quad (2)$$

In [?] they prove that it always admits a unique Nash Equilibrium x^* and has Price of Anarchy $9/8$ if G is undirected and $O(n)$ in the directed case. We denote an instance of this game as $I = (G, s, a)$. The FJ-model can be viewed as the best-response dynamics of the of the opinion formation game. More precisely, each agent i has an opinion $x_1(0), \dots, x_n(0)$ and at round $t \geq 1$, suffers the cost (3) and updates her opinion so as to minimize her individual cost

$$x_i(t) = \operatorname{argmin}_{x \in [0,1]} (1 - a_i) \sum_{j \in N_i} (x - x_j(t-1))^2 + a_i (x - s_i)^2 \quad (3)$$

In [GS14] it is proved that for any instance $I = (G, s, a)$ the opinion vector $x(t) = (x_1(t), \dots, x_n(t))$ converges to the unique equilibrium point x^* . Moreover, this is achieved by an update rule that is very simple but more importantly it is *rational*, in the sense that each agent i adopts this rule in order to minimize her individual cost.

1.2 Opinion Formation Games with Random Payoffs

In the game defined by (2), agent's i cost $C_i(x_i, x_{-i})$ is a deterministic function of the opinion vector x . Many recent works (see e.g. [ZLZ17]) study games with random payoffs, that is agent's i cost ($C_i(x_i, x_{-i})$) is a random variable. The random payoff setting can be much more realistic, since randomness may naturally occur because of incomplete information, noise or other stochastic factors. Motivated by this line of research we introduce a random payoff variant of the opinion formation game (2).

Definition 1. Let $I = (G, s, a)$ an instance of the opinion formation game and x the opinion vector. Each agent i ,

- picks uniformly at random one of her neighbors $j \in N_i$
- suffers cost $C_i(x_i, x_{-i}) = (1 - a_i) \sum_{j \in N_i} (x_i - x_j)^2 + a_i (x_i - s_i)^2$

This game is more compatible to a realistic settings since in real world networks (e.g. Facebook, Twitter e.t.c.), each agent may have several hundreds of friends. As a result is far more reasonable to assume that every day, each agent meets a random small subset of her acquaintances and suffers a cost based on how much she disagrees with them. Since the expected cost of our random payoff variant equals the cost of game (2), $\mathbb{E}[C(x_i, x_{-i})] = (1 - a_i) \sum_{j \in N_i} (x_i - x_j)^2 + a_i (x_i - s_i)^2$, the results about the existence of a unique equilibrium, and the PoA bounds also hold in our variant. However, the best-response dynamics in the random payoff game does not converge to the equilibrium. In this work, we investigate whether there exists natural dynamics that leads the system to the equilibrium point x^* .

1.3 Our Results and Techniques

In this work we propose Algorithm 1 as a dynamics for the random payoff game of Definition 1. Note that the dynamics described in Algorithm 1 is the *fictitious play* in the game defined by the instance $I = (G, s, a)$. Generally speaking *fictitious play* does not guarantee convergence to the equilibrium. In Section 3 we show

Algorithm 1 Fictitious Play

- 1: Initially, each agent i has an opinion $x_1(0), \dots, x_n(0)$.
- 2: At round $t \geq 1$, each agent i meets uniformly at random one of her friends $j \in N_i$
- 3: Suffers cost $C_i^t(x_i(t-1), x_j(t-1)) = (1 - a_i)(x_i(t-1) - x_j(t-1))^2 + a_i(x_i(t-1) - s_i)^2$
- 4: Updates opinion

$$x_i(t) = \operatorname{argmin}_{x \in [0,1]} \sum_{\tau=1}^t C_i^\tau(x, x_{\pi_i(\tau)}(\tau-1)) \quad (4)$$

that, for our game with random payoffs, fictitious play converges to equilibrium, with the following rate.

Theorem 1. *Let $I = (G, s, \alpha)$ be any instance of the opinion formation game of Definition 1 with equilibrium $x^* \in [0, 1]^n$. The opinion vector x^t produced by Algorithm 1 after t rounds satisfies*

$$\mathbb{E} [\|x^t - x^*\|_\infty] \leq C \sqrt{\log n} \frac{(\log t)^2}{t^{\min(1/2, \rho)}},$$

where $\rho = \min_{i \in V} a_i$ and C is a universal constant.

The update rule (4) guarantees convergence while vastly reducing the information exchange between the agents at each round. Namely, each agent learns the opinion of only one of her neighbors whereas in the classical FJ-model (1) each agent learns the opinions of all of her neighbors. In terms of total communication needed to get within distance ε of the equilibrium, the update rule (4) needs $O(|V| \log |V|)$ communication while (1) needs $O(|E|)$. Of course for this difference to be significant we need each agent to have at least $O(\log |V|)$ neighbors. A large social network like Facebook has approximately 2 billion users and each user has usually more than 100 friends which far more than $\log(2 \cdot 10^9)$.

Apart from converging to the equilibrium, our update rule (4) is also a rational behavioral assumption since it ensures *no-regret* for the agents. Having no-regret means that the average cost for each agent i after T rounds is close to the average cost that she would suffer by expressing any fixed opinion. This is a very important feature of our update rule because even players that selfishly will to minimize their incurred cost, could choose to play according to it. In Section 4 we show the following theorem.

Theorem 2. *For every instance $I = (G, s, a)$, for every agent i*

$$\sum_{\tau=1}^t C_i^\tau(x_i(\tau)) \leq \min_{x \in [0,1]} \sum_{\tau=1}^t C_i^\tau(x) + O(\log t)$$

Even though our update rule (4) has the above desired properties it only achieves convergence rate of $\tilde{O}(1/\sqrt{t})$ for a fixed instance I with $\rho > 1/2$. For such an instance the original FJ-model outperforms our update rule since it achieves convergence rate $O(1/2^t)$. We investigate whether this gap is due to an inherent characteristic of the random payoff setting i.e. learning the opinion of just one random neighbor vs learning all of them or we could find another natural dynamics that converges exponentially fast. Can the agents select another no-regret algorithm that guarantees an asymptotically faster convergence rate? In Section 4 we investigate the following question

Question. *Is there a no-regret algorithm that the agents can choose such that for any instance $I = (G, s, a)$, $\mathbb{E} [\|x_t - x^*\|_\infty] = O(1/t)$?*

To answer this question we first show that the existence of a no-regret algorithm (that achieves this convergence rate to the equilibrium) implies the existence of an algorithm that uses i.i.d samples from a Bernoulli random variable $B(p)$ to estimate its success probability p with the same asymptotic error rate. We use sample complexity lower bound techniques from statistics to show that such a no-regret algorithm is unlikely to exist. We show that updating the opinion according to no-regret algorithms leads to update rules that are totally ignorant of the specific graph structure G . Perhaps surprisingly, in Section 8 we present an update rule that takes into account only the maximum degree of the graph and for any fixed instance I achieves convergence rate $O(1/2^{\sqrt{t}})$.

1.4 Related Work

2 Preliminaries

2.1 Notation

Let $G(V, E)$ be a graph. We denote by V the set of agents, and $N_i \subset V$ the set of neighbors of agent i . Let $n = |V|$. We denote by $x(t) \in [0, 1]^n$ the vector of opinions of the agents at round t and $x_i(t)$ its i -th coordinate. We denote by $Q([0, 1])$ the set of rationals in $[0, 1]$.

2.2 Online Convex Optimization and No Regret Algorithms

The *Online Convex Optimization* (OCO) framework can be seen as a game played between a player and an adversary. Let $K \subseteq \mathbf{R}^n$ be a convex set and a set of functions \mathcal{F} defined over K . Both player and adversary know the sets K and \mathcal{F} . At round t ,

1. the agent chooses $x_t \in K$.
2. the adversary observes the x_t and selects a function $f_t(x) \in \mathcal{F}$.
3. the player receives cost $f_t(x_t)$.

The goal of the player is to pick x_t based on the history (f_1, \dots, f_{t-1}) in a way that minimizes the total cost. We emphasize that the agent has to choose x_t before seeing f_t , otherwise the problem becomes trivial.

Definition 2. A for OCO is a sequence of functions, $\{A_t\}_{t=1}^\infty$ where $A_t : \mathcal{F}^{t-1} \mapsto K$. The regret of the algorithm A as round T is

$$R_A(T) = \sup_{\{f_1, \dots, f_T\} \subseteq \mathcal{F}^T} \left(\sum_{t=1}^T f_t(x^t) - \min_{x \in K} \sum_{t=1}^T f_t(x) \right),$$

where $x^t = A^t(f_1, \dots, f_{t-1})$. An algorithm A is no-regret if $R_A(T) = o(T)$.

In our problem agent i selects at round t an opinion $x_i(t) \in [0, 1]$ and then suffers a cost $C_i^t(x_i(t))$. Obviously in our case $K = [0, 1]$ and $\mathcal{F} = \{x \mapsto a_i(x - s_i)^2 + (1 - a_i)(x - c)^2 : a_i, s_i, c \in [0, 1]\}$.

Definition 3 (no-regret Dynamics). Let A be a no-regret algorithm for \mathcal{F}, K . For an instance $I = (G, s, a)$ of the opinion formation game given in Definition 1 we define the following no-regret dynamics.

- Initially, each agent i has an opinion $x_1(0), \dots, x_n(0)$.
- At round $t \geq 1$, each agent i :
 - Meets uniformly at random one of her friends $j \in N_i$.
 - Suffers cost $C_i^t(x_i(t-1), x_j(t-1)) = (1 - a_i)(x_i(t-1) - x_j(t-1))^2 + a_i(x_i(t-1) - s_i)^2$.
 - Updates her opinion using A , $x_i(t) = A_t(C_i^1, \dots, C_i^{t-1})$.

2.3 Estimating Bernoulli distributions

Assume that we have access to the random variables Y_1, \dots, Y_t , that are mutually independent and each Y_i is distributed according to a Bernoulli distribution with $p \in [0, 1]$, $X_i \sim B(p)$. Initially the parameter p is unknown. The task is to observe the realization of y_1, \dots, y_t and output an estimate $\hat{p} = \theta_t(y_1, \dots, y_t)$ according to a function $\hat{\theta}_t : \{0, 1\}^t \mapsto [0, 1]$, that is close to the unknown parameter p .

Definition 4. An estimator $\hat{\theta}$ is a collection of functions, $\{\hat{\theta}_t\}_{t=1}^\infty$, where $\hat{\theta}_t : \{0, 1\}^t \mapsto [0, 1]$.

A well known example of such an estimator is the *sample mean*, where $\hat{\theta}_t = \frac{\sum_{i=1}^t Y_i}{t}$. Obviously for any efficient estimator, the estimate \hat{p} converges to the unknown parameter p as t grows. The following definition is the standard metric for the efficiency of an estimator $\hat{\theta}$.

Definition 5. For an estimator $\hat{\theta} = \{\hat{\theta}_t\}_{t=1}^\infty$ we define its error rate $R_p(t) = E_p[|\hat{\theta}_t(Y_1, \dots, Y_t) - p|]$

$$\text{where } E_p[|\hat{\theta}_t(Y_1, \dots, Y_t) - p|] = \sum_{(y_1, \dots, y^t) \in \{0, 1\}^t} |\hat{\theta}_t(y^1, \dots, y^t) - p| \cdot p^{\sum_{i=1}^t y_i} (1-p)^{t-\sum_{i=1}^t y_i}$$

The quantity $E_p[|\hat{\theta}_t(Y_1, \dots, Y_t) - p|]$ is the expected distance of the estimated value $\hat{\theta}_t$ from the parameter p , when the distribution of the samples is $B(p)$. To simplify notation we also denote it as $E_p[|\hat{\theta}_t - p|]$. The error rate $R_p(t)$ quantifies the rate of convergence of the estimated value $\hat{p} = \theta_t(Y_1, \dots, Y_t)$ to the real parameter p .

Since p is unknown, any meaningful estimator \hat{p} must guarantee that for all $p \in [0, 1]$, $\lim_{t \rightarrow \infty} R_p(t) = 0$. For example the *sample mean* has error rate $R_p(t) \leq \frac{1}{2\sqrt{t}}$ for any $p \in [0, 1]$ and clearly satisfies the above requirement. In section 4 we investigate the following question,

Question. Is there an estimator $\hat{\theta}$ such that for all $p \in [0, 1]$, $\lim_{t \rightarrow \infty} tR_p(t) = 0$?

3 Fictitious Play Convergence Rate

In this Section we prove that the fictitious play described as Algorithm 1 converges to the equilibrium. We denote by V the set of agents, and $N_i \subset V$ the set of neighbors of agent i . Let $x(t) \in [0, 1]^n$ be the opinion vector at round t . We denote by $x_i(t)$ the opinion of agent i at round t . At round $t + 1$, $x(t + 1)$ is constructed in the following way. At first, each agent i gets to see the opinion $x_j(t)$, where $j \in N_i$ is picked uniformly at random from the set of her neighbors. Let W_i^t be the random variable corresponding to the selected neighbor. Now each agent $i \in V$ updates her opinion as follows

$$x_i(t + 1) = (1 - \alpha_i) \frac{\sum_{\tau=1}^{t+1} x_{W_i^\tau}(\tau - 1)}{t + 1} + \alpha_i s_i.$$

Let $x^* \in [0, 1]^n$ be the unique equilibrium point of the given instance I . We prove that the above stochastic process has the following convergence rate to x^* .

Theorem 3. Let $I = (G, s, \alpha)$ be any instance of the opinion formation game of Definition 1 with equilibrium $x^* \in [0, 1]^n$. The opinion vector x^t produced by Algorithm 1 after t rounds satisfies

$$E[\|x^t - x^*\|_\infty] \leq C \sqrt{\log n} \frac{(\log t)^2}{t^{\min(1/2, \rho)}},$$

where $\rho = \min_{i \in V} \alpha_i$ and C is a universal constant.

Proof. Setting $p = 1/\sqrt{t}$ in Lemma 3 yields the result. ■

We start by stating the standard Hoeffding bound

Lemma 1 (Hoeffding's Inequality). Let $X = (X_1 + \dots + X_t)/t$, $X_i \in [0, 1]$. Then, $\mathbf{P}[|X - \mathbf{E}[X]| > \lambda] < 2e^{-2t\lambda^2}$.

Lemma 2. With probability at least $1 - p$, $\|x^t - x^*\|_\infty \leq e(t)$, where $e(t)$ satisfies the following recursive relation

$$e(t) = \delta(t) + (1 - \alpha) \frac{\sum_{\tau=0}^{t-1} e(\tau)}{t} \text{ and } e(0) = \|x^0 - x^*\|_\infty,$$

where $\delta(t) = \sqrt{\frac{\log(\pi^2 nt^2 / (6p))}{t}}$.

Corollary 1. The function $e(t)$ satisfies the following recursive relation

$$e(t+1) - e(t) + \alpha \frac{e(t)}{t+1} = \delta(t+1) - \delta(t) + \frac{\delta(t)}{t+1}$$

Proof. As we have already mentioned for any instance I there exists a unique equilibrium vector x^* . Since $W_i^\tau \sim \mathcal{U}(N_i)$ we have that $\mathbf{E}[x_{W_i^\tau}^*] = \frac{\sum_{j \in N_i} x_j^*}{|N_i|}$. Since W_i^τ are independent random variables, we can use Hoeffding's inequality (Lemma 1) to get

$$\mathbf{P} \left[\left| \frac{\sum_{\tau=1}^t x_{W_i^\tau}^*}{t} - \frac{\sum_{j \in N_i} x_j^*}{|N_i|} \right| > \delta(t) \right] < \frac{6}{\pi^2} \frac{p}{nt^2},$$

where $\delta(t) = \sqrt{\frac{\log(\pi^2 nt^2 / (6p))}{t}}$. Therefore, by the union bound

$$\begin{aligned} & \mathbf{P} \left[\text{for all } t \geq 1 : \max_{i \in V} \left| \frac{\sum_{\tau=1}^t x_{W_i^\tau}^*}{t} - \frac{\sum_{j \in N_i} x_j^*}{|N_i|} \right| > \delta(t) \right] \leq \\ & \sum_{t=1}^{\infty} \mathbf{P} \left[\max_{i \in V} \left| \frac{\sum_{\tau=1}^t x_{W_i^\tau}^*}{t} - \frac{\sum_{j \in N_i} x_j^*}{|N_i|} \right| > \delta(t) \right] \leq \\ & \sum_{i=1}^{\infty} \frac{6}{\pi^2} \frac{1}{t^2} \sum_{i=1}^n \frac{p}{n} = p \end{aligned}$$

As a result with probability $1 - p$ we have that for all t

$$\left| \frac{\sum_{\tau=1}^t x_{W_i^\tau}^*}{t} - \frac{\sum_{j \in N_i} x_j^*}{|N_i|} \right| \leq \delta(t) \quad (5)$$

We will our claim by induction. We assume that $\|x^\tau - x^*\|_\infty \leq e(\tau)$ for all $\tau \leq t-1$. Then

$$\begin{aligned} x_i(t) &= (1 - \alpha_i) \frac{\sum_{\tau=1}^t x_{W_i^\tau}^*(\tau-1)}{t} + \alpha_i s_i \\ &\leq (1 - \alpha_i) \frac{\sum_{\tau=1}^t x_{W_i^\tau}^* + \sum_{\tau=1}^t e(\tau-1)}{t} + \alpha_i s_i \end{aligned} \quad (6)$$

$$\begin{aligned} &\leq (1 - \alpha_i) \left(\frac{\sum_{\tau=1}^t x_{W_i^\tau}^*}{t} + \frac{\sum_{\tau=0}^{t-1} e(\tau)}{t} \right) + \alpha_i s_i \\ &\leq (1 - \alpha_i) \left(\frac{\sum_{j \in N_i} x_j^*}{|N_i|} + \delta(t) + \frac{\sum_{\tau=0}^{t-1} e(\tau)}{t} \right) + \alpha_i s_i \end{aligned} \quad (7)$$

$$\leq x_i^* + \delta(t) + (1 - \alpha) \left(\frac{\sum_{\tau=0}^{t-1} e(\tau)}{t} \right)$$

We get (6) from the induction step and (7) from inequality (5). Similarly, we can prove that $x_i(t) \geq x_i^* - \delta(t) - (1 - \alpha) \frac{\sum_{\tau=1}^t e(\tau)}{t}$. As a result $\|x_i(t) - x^*\|_\infty \leq e(t)$. ■

In order to bound the convergence time of the system, we just need to bound the convergence rate of the function $e(t)$. The following lemma provides us with a simple upper bound for the convergence rate of our process.

Lemma 3. *Let $e(t)$ be a function satisfying the recursion of Corollary 1. Then with probability at least $1 - p$ we have that*

$$e(t) = \begin{cases} O\left(\sqrt{\log\left(\frac{n}{p}\right)} \frac{(\log t)^{3/2}}{t^\alpha}\right) & \text{if } \alpha \leq 1/2 \\ O\left(\sqrt{\log\left(\frac{n}{p}\right)} \frac{(\log t)^{3/2}}{t^{1/2}}\right) & \text{if } \alpha > 1/2 \end{cases}$$

Proof. At first, since $\frac{\delta(t)}{t}$ is a decreasing function for $p \leq 1/4$, we have that $e(t) \leq (1 - \frac{\alpha}{t})e(t-1) + g(t)$, where $g(t) = \frac{\delta(t)}{t}$.

$$\begin{aligned} e(t) &\leq (1 - \frac{\alpha}{t})e(t-1) + g(t) \\ &\leq (1 - \frac{\alpha}{t})(1 - \frac{\alpha}{t-1})e(t-2) + (1 - \frac{\alpha}{t})g(t-1) + g(t) \\ &\leq (1 - \frac{\alpha}{t}) \cdots (1 - \alpha)e(0) + \sum_{\tau=1}^t g(\tau) \prod_{i=\tau+1}^t (1 - \frac{\alpha}{i}) \\ &\leq \frac{e(0)}{t^\alpha} + \sum_{\tau=1}^t g(\tau) e^{-\alpha \sum_{i=\tau+1}^t \frac{1}{i}} \\ &\leq \frac{e(0)}{t^\alpha} + \sum_{\tau=1}^t g(\tau) e^{-\alpha(H_t - H_\tau)} \\ &\leq \frac{e(0)}{t^\alpha} + e^{-\alpha H_t} \sum_{\tau=1}^t g(\tau) e^{\alpha H_\tau} \\ &\leq \frac{e(0)}{t^\alpha} + \frac{O\left(\sqrt{\log\left(\frac{n}{p}\right)}\right)}{t^\alpha} \sum_{\tau=1}^t \tau^\alpha \frac{\sqrt{\log \tau}}{\tau^{3/2}} \end{aligned}$$

We observe that

$$\sum_{\tau=1}^t \tau^\alpha \frac{\sqrt{\log \tau}}{\tau^{3/2}} \leq \int_{\tau=1}^t \tau^\alpha \frac{\sqrt{\log \tau}}{\tau^{3/2}} d\tau,$$

since $\tau^\alpha \frac{\sqrt{\log \tau}}{\tau^{3/2}}$

- If $\alpha \leq 1/2$ then

$$\int_{\tau=1}^t \tau^\alpha \frac{\sqrt{\log \tau}}{\tau^{3/2}} d\tau \leq \int_{\tau=1}^t \tau^{1/2} \frac{\sqrt{\log \tau}}{\tau^{3/2}} d\tau = O((\log t)^{3/2})$$

- If $\alpha > 1/2$ then

$$\begin{aligned} \int_{\tau=1}^t \tau^\alpha \frac{\sqrt{\log \tau}}{\tau^{3/2}} d\tau &= \int_{\tau=1}^t \tau^{\alpha-1/2} \frac{\sqrt{\log \tau}}{\tau} d\tau \\ &= \frac{2}{3} \int_{\tau=1}^t \tau^{\alpha-1/2} ((\log \tau)^{3/2})' d\tau \\ &= \frac{2}{3} (\log t)^{3/2} - (\alpha - 1/2) \frac{2}{3} \int_{\tau=1}^t \tau^{\alpha-3/2} (\log \tau)^{3/2} d\tau \\ &= O((\log t)^{3/2}) \end{aligned}$$

■

4 Fictitious Play is no-regret

We consider the following online convex optimization problem. At each time step t , the player i selects a real number x^t and a function $f^t(x)$ arrives. The player then suffers $f^t(x^t)$ cost. The functions $f^t(x)$ have the following form:

$$f^t(x) = \alpha(x - s_i)^2 + (1 - \alpha)(x - a_t)^2$$

where $s_i, \alpha \in [0, 1]$ and are independent of t and $a_t \in [0, 1]$. In other words, the function f^t is uniquely determined by the number a_t . We show that for this class of functions *fictitious play* admits no regret.

Theorem 4. Let f^t be a sequence of functions, where each function has the form: $f^t(x) = \alpha(x - s_i)^2 + (1 - \alpha)(x - a_t)^2$, where $s_i, \alpha \in [0, 1]$ and are independent of t and $a_t \in [0, 1]$. If we define $x^t = \operatorname{argmin}_{x \in [0, 1]} \sum_{\tau=1}^{t-1} f^\tau(x)$ then for each time T

$$\sum_{t=1}^T f^t(x^t) \leq \min_{x \in [0, 1]} \sum_{t=1}^T f^t(x) + O(\log T)$$

In order to prove this claim, we will first define a similar rule y^t that takes into account the function f^t for the "prediction" at time t . Intuitively, this guarantees that the rule admits no regret.

Lemma 4. Let $y^t = \operatorname{argmin}_{x \in [0, 1]} \sum_{\tau=1}^t f^\tau(x)$ then

$$\sum_{t=1}^T f^t(y^t) \leq \min_{x \in [0, 1]} \sum_{t=1}^T f^t(x)$$

Proof. By definition of y^t , $\min_{x \in [0, 1]} \sum_{t=1}^T f^t(x) = \sum_{t=1}^T f^t(y^T)$, so

$$\begin{aligned} \sum_{t=1}^T f^t(y^t) - \min_{x \in [0, 1]} \sum_{t=1}^T f^t(x) &= \sum_{t=1}^T f^t(y^t) - \sum_{t=1}^T f^t(y^T) \\ &= \sum_{t=1}^{T-1} f^t(y^t) - \sum_{t=1}^{T-1} f^t(y^T) \\ &\leq \sum_{t=1}^{T-1} f^t(y^t) - \sum_{t=1}^{T-1} f^t(y^{T-1}) \\ &= \sum_{t=1}^{T-2} f^t(y^t) - \sum_{t=1}^{T-2} f^t(y^{T-1}) \end{aligned}$$

Continuing in the same way, we get $\sum_{t=1}^T f^t(y^t) \leq \min_{x \in [0, 1]} \sum_{t=1}^T f^t(x)$. ■

Now we can derive some intuition for the reason that *fictitious play* admits no regret. Since the cost incurred by the sequence y^t is at most that of the best fixed strategy, we can compare the cost incurred by x^t with that of y^t . However, for each t the numbers x^t and y^t are quite close and as a result the difference in their cost must be quite small.

Lemma 5. For all t , $f^t(x^t) \leq f^t(y^t) + 2\frac{1-\alpha}{t} + \frac{(1-\alpha)^2}{t^2}$.

Proof. We first prove that for all t ,

$$|x^t - y^t| \leq \frac{1-\alpha}{t}. \tag{8}$$

By definition $x^t = \alpha s_i + (1 - \alpha) \frac{\sum_{\tau=1}^{t-1} a_\tau}{t-1}$ and $y^t = \alpha s_i + (1 - \alpha) \frac{\sum_{\tau=1}^t a_\tau}{t}$.

$$\begin{aligned} |x^t - y^t| &= (1 - \alpha) \left| \frac{\sum_{\tau=1}^{t-1} a_\tau}{t-1} - \frac{\sum_{\tau=1}^t a_\tau}{t} \right| \\ &= (1 - \alpha) \left| \frac{\sum_{\tau=1}^{t-1} a_\tau - (t-1)a_t}{t(t-1)} \right| \\ &\leq \frac{1 - \alpha}{t} \end{aligned}$$

The last inequality follows from the fact that $a_\tau \in [0, 1]$. We now use inequality (8) to bound the difference $f^t(x^t) - f^t(y^t)$.

$$\begin{aligned} f^t(x^t) &= \alpha(x^t - s_i)^2 + (1 - \alpha)(x^t - a_t)^2 \\ &\leq \alpha(y^t - s_i)^2 + 2\alpha |y^t - s_i| |x^t - y^t| + \alpha |x^t - y^t|^2 \\ &\quad + (1 - \alpha)(y^t - a_t)^2 + 2(1 - \alpha) |y^t - a_t| |x^t - y^t| + (1 - \alpha) |x^t - y^t|^2 \\ &\leq f^t(y^t) + 2 |x^t - y^t| + |y^t - x^t|^2 \\ &\leq f^t(y^t) + 2 \frac{1 - \alpha}{t} + \frac{(1 - \alpha)^2}{t^2} \end{aligned}$$

■

Theorem 4 easily follows since:

$$\begin{aligned} \sum_{t=1}^T f^t(x^t) &\leq \sum_{t=1}^T f^t(y^t) + \sum_{t=1}^T 2 \frac{1 - \alpha}{t} + \sum_{t=1}^T \frac{(1 - \alpha)^2}{t^2} \\ &\leq \min_{x \in [0, 1]} \sum_{t=1}^T f^t(x) + 2(1 - \alpha)(\log T + 1) + (1 - \alpha) \frac{\pi^2}{6} \\ &\leq \min_{x \in [0, 1]} \sum_{t=1}^T f^t(x) + O(\log T) \end{aligned}$$

5 Lower Bound

Theorem 5. Suppose $\hat{\theta}^t$ is an estimator for the probability p of a Bernoulli random variable that takes t samples. Then, for every interval $[a, b]$ which is contained in $[0, 1]$ and for every $c > 0$, it holds that:

$$\lim_{t \rightarrow \infty} t^{1+c} \int_a^b \mathbf{E} [|\hat{\theta}^t - p|_\infty] dp = \infty$$

In order to prove this, we need the following result:

Lemma 6 (Fano's Inequality). Suppose we have an estimator $\hat{\theta}^t$ for the probability of a Bernoulli distribution. The unknown distribution belongs to a family of distributions \mathcal{P} with the property that the probabilities of every two distributions in this family differ by at least 2δ , where δ is a positive real number. Then, for every subset P_1, P_2, \dots, P_n of distributions from this family, where p_i is the probability of distribution P_i , the following inequality holds:

$$\frac{1}{n} \sum_{i=1}^n \mathbf{E} [|\hat{\theta}^t - p_i|] \geq \delta \left(1 - \frac{I(X; V) + \log 2}{\log n} \right)$$

We are now going to use Theorem (6) to prove the following lemma, which provides a lower bound for the mean of the expectation of errors of the estimator.

Lemma 7. Suppose we choose t Bernoulli distributions $\{P_i\}_{i=0}^{t-1}$ with the following probabilities:

$$p_i = a + \frac{b-a}{2t} + i \frac{b-a}{t}, i = 0, \dots, t-1$$

. Then:

$$\frac{1}{t} \sum_{i=1}^t \mathbf{E} [|\hat{\theta}^t - p_i|] \geq \frac{b-a}{2t} \left(c_1 - \frac{c_2}{t} \right)$$

where $c_1, c_2 > 0$ and $c_1 > \frac{1}{2}$.

Proof. Without loss of generality, we can assume that t is a multiple of 5. For simplicity, we set $E_i = \mathbf{E} [|\hat{\theta}^t - p_i|]$. By the choice of p_i , we have that

$$|p_i - p_j| \geq \frac{b-a}{t}$$

, for every distinct i, j in the family. This means that for the family of t distributions that we picked, the property of Lemma (6) is satisfied for $\delta = \frac{b-a}{2t}$. Thus, by dividing the t distributions into groups of 5 and applying Lemma (6) to each one, we obtain:

$$\begin{aligned} \frac{\sum k = i^{i+4}}{5} &\geq \frac{b-a}{2t} \left(1 - \frac{I(X; V) + \log 2}{\log n} \right) \\ &= \frac{b-a}{2t} \left(1 - \frac{\log 2}{\log 5} - \frac{I(X; V)}{\log 5} \right) \\ &= \frac{b-a}{2t} \left(c_1 - \frac{I(X; V)}{\log 5} \right) \end{aligned} \tag{9}$$

where $c_1 = 1 - \frac{\log 2}{\log 5} > \frac{1}{2}$. We are now going to upper bound the mutual information $I(X; V)$ for every group of 5 distributions. We denote by P_i^t the distribution on vectors with t coordinates, where each coordinate is sampled independently from P_i . Let $U = \{i, i+1, i+2, i+3, i+4\}$ be a family of 5 distributions. Using a well known inequality REFERENCE NEEDED HERE:

$$\begin{aligned} I(X; V) &\leq \frac{1}{5^2} \sum_{\substack{i, j \in U \\ i \neq j}} D_{kl} (P_i^t, P_j^t) \\ &\leq D_{kl} (P_i^t, P_{i+4}^t) \\ &= t D_{kl} (P_i, P_{i+4}) \\ &\leq t \frac{(p_i - p_{i+4})^2}{p_i(1-p_i)} \\ &\leq t \left(\frac{1}{a(1-a)} + \frac{1}{b(1-b)} \right) (p_i - p_{i+4})^2 \\ &= t \left(\frac{1}{a(1-a)} + \frac{1}{b(1-b)} \right) \left(\frac{b-a}{2t} + i \frac{b-a}{t} - \frac{b-a}{t} - (i+4) \frac{b-a}{t} \right)^2 \\ &= t \left(\frac{1}{a(1-a)} + \frac{1}{b(1-b)} \right) \frac{16(b-a)^2}{t^2} \\ &= \frac{16(b-a)^2}{t} \left(\frac{1}{a(1-a)} + \frac{1}{b(1-b)} \right) \end{aligned}$$

So, if we set

$$c_2 = 16 \frac{(b-a)^2}{\log 5} \left(\frac{1}{a(1-a)} + \frac{1}{b(1-b)} \right)$$

then by equation (9) we obtain:

$$\sum_{k=i}^{i+4} \geq \frac{5(b-a)}{2t} \left(c_1 - \frac{c_2}{t} \right) \quad (10)$$

Inequality (10) holds for every group of 5 distributions. By summing up for all groups:

$$\begin{aligned} \sum_{i=0}^{t-1} E_i &\geq t \frac{b-a}{2t} \left(c_1 - \frac{c_2}{t} \right) \\ &= \frac{b-a}{2} \left(c_1 - \frac{c_2}{t} \right) \end{aligned}$$

which is what we wanted to prove. ■

Next, we state a lemma regarding an upper bound on the expectation of the error in each p_i .

Lemma 8. For every $p \in [-(b-a)/2t + p_i, p_i + (b-a)/2t]$:

$$\mathbf{E} [|\hat{\theta}^t - p_i|]_{p_i} \leq c_a^b \mathbf{E} [|\hat{\theta}^t - p|]_p + |p - p_i|$$

where c_a^b is a constant that depends on a, b .

Now, we find an upper bound for the quantity $\sum_{i=0}^{t-1} E_i$. Our goal is to find an upper bound where the coefficient of $(b-a)/t$ is less than $1/4$.

Lemma 9.

$$\frac{1}{t} \sum_{i=0}^{t-1} E_i \leq \frac{c_a^b}{b-a} \int_a^b \mathbf{E} [|\hat{\theta}^t - p|]_p dp + \frac{b-a}{4t}$$

Proof. By Lemma (8) we get that for every $p \in [-(b-a)/2t + p_i, p_i + (b-a)/2t]$:

$$E_i \leq c_a^b \mathbf{E} [|\hat{\theta}^t - p|]_p + |p_i - p|$$

We integrate both sides of the equation with respect to p and get:

$$\int_{p_i - \frac{b-a}{2t}}^{p_i + \frac{b-a}{2t}} E_i dp \leq c_a^b \int_{p_i - \frac{b-a}{2t}}^{p_i + \frac{b-a}{2t}} \mathbf{E} [|\hat{\theta}^t - p|]_p dp + \int_{p_i - \frac{b-a}{2t}}^{p_i + \frac{b-a}{2t}} |p_i - p| dp$$

Hence:

$$\begin{aligned} \frac{b-a}{t} E_i &\leq c_a^b \int_{p_i - \frac{b-a}{2t}}^{p_i + \frac{b-a}{2t}} \mathbf{E} [|\hat{\theta}^t - p|]_p dp + 2 \int_{p_i - \frac{b-a}{2t}}^{p_i} |p_i - p| dp \\ &= c_a^b \int_{p_i - \frac{b-a}{2t}}^{p_i + \frac{b-a}{2t}} \mathbf{E} [|\hat{\theta}^t - p|]_p dp + \frac{1}{4} \frac{(b-a)^2}{t^2} \end{aligned}$$

Therefore:

$$\frac{E_i}{t} \leq \frac{c_a^b}{b-a} \int_{p_i - \frac{b-a}{2t}}^{p_i + \frac{b-a}{2t}} \mathbf{E} [|\hat{\theta}^t - p|]_p dp + \frac{b-a}{4t^2} \quad (11)$$

By summing up Equation (11) for all t intervals, we get:

$$\begin{aligned} \frac{\sum_{i=0}^{t-1} E_i}{t} &\leq \frac{c_a^b}{b-a} \sum_{i=0}^{t-1} \int_{p_i - \frac{b-a}{2t}}^{p_i + \frac{b-a}{2t}} \mathbf{E} [|\hat{\theta}^t - p|]_p dp + \frac{b-a}{4t} \\ &= \frac{c_a^b}{b-a} \int_a^b \mathbf{E} [|\hat{\theta}^t - p|]_p dp + \frac{b-a}{4t} \end{aligned}$$
■

Now, by combining Lemmas (7) and (9) we get:

$$\frac{c_a^b}{b-a} \int_a^b \mathbf{E} [|\hat{\theta}^t - p|]_p dp \geq \frac{b-a}{2t} \left(\left(c_1 - \frac{1}{2} \right) - \frac{c_2}{t} \right)$$

By multiplying by t^{1+c} we get:

$$\frac{c_a^b}{b-a} t^{1+c} \int_a^b \mathbf{E} [|\hat{\theta}^t - p|]_p dp \geq \frac{b-a}{2} t^c \left(\left(c_1 - \frac{1}{2} \right) - \frac{c_2}{t} \right) \quad (12)$$

The coefficient of t^c in the right hand side of (12) is positive, so by sending $t \rightarrow \infty$ we get:

$$\lim_{t \rightarrow \infty} t^{1+c} \int_a^b \mathbf{E} [|\hat{\theta}^t - p|] dp = \infty$$

6 Lower bound

Theorem 6. Let P be a protocol run by each agent that produces vector output x^t after t iterations. Suppose that for every graph G and for every choice of parameters $s_i, \alpha_i \in [0, 1]$ the error of the output is:

$$\mathbf{E} [\|x^t - x^*\|_\infty] = f(n, t)$$

where n is the number of nodes in the graph. Then,

$$\lim_{t \rightarrow \infty} \sqrt{t} f(t, t) > 0$$

Proof. Let G be a graph that consists of 2 star topologies. That is, we have 2 center nodes v_1, v_2 . Each of the two center nodes has exactly t neighbours, which have degree 1. We will call these nodes the leaves of the graph. This graph has $2t + 2$ nodes. We set $\alpha_i = 1$ for all the leaves and $\alpha_i = 1/2$ for the two centers. For exactly $t/2 + \sqrt{t}/2$ leaves of v_1 we set $s_i = 1$, for the rest $s_i = 0$. Similarly, for exactly $t/2 - \sqrt{t}/2$ leaves of v_2 we set $s_i = 1$, for the rest $s_i = 0$. By direct calculation, we get that the equilibrium point for this graph is the following:

$$x_i^* = \begin{cases} s_i & \text{if } i \text{ is a leaf node} \\ \frac{1}{4} + \frac{1}{4\sqrt{t}} & \text{if } i = v_1 \\ \frac{1}{4} - \frac{1}{4\sqrt{t}} & \text{if } i = v_2 \end{cases}$$

Since the graph has $2t + 2$ nodes:

$$\mathbf{E} [\|x^t - x^*\|_\infty] = f(2t + 2, t) \quad (13)$$

On the other hand, since $\|x^t - x^*\|_\infty \geq \max\{|x_{v_1}^t - x_{v_1}^*|, |x_{v_2}^t - x_{v_2}^*|\}$ we get:

$$\mathbf{E} [\|x^t - x^*\|_\infty] \geq \frac{1}{2} \mathbf{E} [|x_{v_1}^t - x_{v_1}^*|] + \frac{1}{2} \mathbf{E} [|x_{v_2}^t - x_{v_2}^*|] \quad (14)$$

Suppose we have a star graph with center node c . At each time t , the oracle returns to the center node either the value of a leaf with $s_i = 1$ (e.g. $f(1, 1, t)$) with probability :

$$p = \frac{\# \text{ leaves with } s_i = 1}{\# \text{ leaves}}$$

or the value of a leaf with $s_i = 0$ (e.g. $f(1, 0, t)$) with probability $1 - p$. As a result, knowing the 0-1 choice of the oracle at time step t , one can compute x_c^t , meaning that $\hat{\theta}^t = 2x_c^t$ is an estimator for a Bernoulli random variable with probability p . This obviously holds for every choice of $s_i \in \{0, 1\}$. This means that in our graph, $\hat{\theta}_1 = 2x_{v_1}^t$ is an estimator for the Bernoulli with probability $p_1 = \frac{1}{2} + \frac{1}{2\sqrt{t}}$ and $\hat{\theta}_2 = 2x_{v_2}^t$ is an estimator for the Bernoulli with probability $p_2 = \frac{1}{2} - \frac{1}{2\sqrt{t}}$. To be more precise, since the protocol is executed

in the same way in all nodes, $\hat{\theta}_1$ and $\hat{\theta}_2$ are the outputs of one estimator $\hat{\theta}$, depending on the distribution from which this estimator gets its input. By combining inequalities (13) and (14) we get:

$$f(2t+2, t) \geq \frac{1}{4}E_{p_1}\{|\hat{\theta} - p_1|\} + \frac{1}{4}E_{p_2}\{|\hat{\theta} - p_2|\} \quad (15)$$

We will now try to bound the right hand side of equation (15). We first notice that:

$$\begin{aligned} E_{p_1}|\hat{\theta} - p_1| &\geq \delta \mathbf{E} [\mathbb{1}\{|\hat{\theta} - p_1| > \delta\}] \\ &= \delta \mathbf{P} [|\hat{\theta} - p_1| > \delta] \end{aligned}$$

The same holds for p_2 . It suffices now to provide a lower bound for this probability. In order to do this, we use a standard reduction from estimation problems to testing problems (see Lecam). In a testing problem, there exist two Bernoulli distributions P_1 and P_2 . First, nature selects with equal probability one of them and let $V \in \{1, 2\}$ denote this random choice. Then, we draw a random vector (Y_1, \dots, Y_t) from the t -fold product distribution P_V^t . The goal is to decide whether $V = 1$ or $V = 2$. More precisely, we have to select a suitable function $\psi : \{0, 1\}^t \mapsto \{1, 2\}$ that minimizes $\mathbf{P}[\psi(Y_1, \dots, Y_t) \neq V]$. First, we state a Lemma that provides a lower bound for this probability.

Lemma 10. *For every function $\psi : \{0, 1\}^t \mapsto \{1, 2\}$:*

$$\mathbf{P}[\psi(Y_1, \dots, Y_t) \neq V] \geq \frac{1}{2} (1 - D_{TV}(P_1^t, P_2^t))$$

where t is the number of samples.

We have:

$$\begin{aligned} \frac{1}{4}E_{p_1}\{|\hat{\theta} - p_1|\} + \frac{1}{4}E_{p_2}\{|\hat{\theta} - p_2|\} &= \frac{1}{4}E\{|\hat{\theta} - p_1| | V = 1\} + \frac{1}{4}E\{|\hat{\theta} - p_2| | V = 2\} \\ &\geq \frac{1}{4\sqrt{t}}\mathbf{P}\left[\hat{\theta} > \frac{1}{2} | V = 1\right] + \frac{1}{4\sqrt{t}}\mathbf{P}\left[\hat{\theta} < \frac{1}{2} | V = 2\right] \end{aligned}$$

We define :

$$\psi = \begin{cases} 1 & \hat{\theta} < \frac{1}{2} \\ 2 & \hat{\theta} > \frac{1}{2} \end{cases}$$

As a result,

$$\begin{aligned} \frac{1}{4\sqrt{t}}\mathbf{P}\left[\hat{\theta} > \frac{1}{2} | V = 1\right] + \frac{1}{4\sqrt{t}}\mathbf{P}\left[\hat{\theta} < \frac{1}{2} | V = 2\right] &\geq \frac{1}{2\sqrt{t}}\left(\mathbf{P}\left[\hat{\theta} > \frac{1}{2}, V = 1\right] + \mathbf{P}\left[\hat{\theta} < \frac{1}{2}, V = 2\right]\right) \\ &= \frac{1}{2\sqrt{t}}(\mathbf{P}[\psi = 2, V = 1] + \mathbf{P}[\psi = 1, V = 2]) \\ &= \frac{1}{2\sqrt{t}}(\mathbf{P}[\psi \neq V]) \\ &\geq \frac{1}{4\sqrt{t}}(1 - D_{TV}(P_1^t, P_2^t)) \\ &\geq \frac{1}{24\sqrt{t}} \end{aligned}$$

So

$$\lim_{t \rightarrow \infty} \sqrt{t}f(2t+2, t) > 0$$

■

6.1 Estimating Bernoulli Distributions

The problem definition is the following:

- There exists an unknown parameter $p \in [0, 1]$
- We receive i.i.d. observations X_i , where $X_i \sim B(p)$
- We output an estimate value $\hat{p} = \hat{\theta}_t(X_1, \dots, X_t)$ according to some function $\hat{\theta}_t : \{0, 1\}^t \mapsto [0, 1]$.

The goal is to appropriately select $\hat{\theta}_t$ such that \hat{p} and p are relatively close. For example a very known selection is

$$\hat{\theta}_t(X_1, \dots, X_t) = \frac{\sum_{i=1}^t X_i}{t}$$

which is the sample mean estimator

Definition 6. An estimator $\hat{\theta}$ is a collection of functions $\{\hat{\theta}_t\}_{t=1}^\infty$, where $\hat{\theta}_t : \{0, 1\}^t \mapsto [0, 1]$.

As the number of samples grows we would expect that an efficient estimator $\hat{\theta}$ produces more and more accurate values for p . The following definition provides us the metric to evaluate the quality of an estimator $\hat{\theta}$.

Definition 7. For any estimator $\hat{\theta}_t$ we define

$$E_p[|\hat{\theta}_t(X_1, \dots, X_t) - p|] = \sum_{(x_1, \dots, x^t) \in \{0, 1\}^t} |\hat{\theta}_t(x^1, \dots, x^t) - p| \cdot p^{\sum_{i=1}^t x^i} (1-p)^{t - \sum_{i=1}^t x^i}$$

The quantity $E_p[|\hat{\theta}_t(X_1, \dots, X_t) - p|]$ is the expected distance of the estimated value $\hat{\theta}_t$ from the parameter p , when the distribution of the samples is $B(p)$. To simplify notation we also denote it as $E_p[|\hat{\theta}_t - p|]$.

Notice that $E_p[|\hat{\theta}_t - p|]$ does not serve as a good metric since an efficient estimator $\hat{\theta}$ must guarantee good estimates for any parameter p . Consider an estimator $\hat{\theta}$ that outputs $1/2$ for all t . In this case, $E_p[|\hat{\theta}_t - p|] = |p - 1/2|$ meaning that this estimator performs optimally when $p = 1/2$, but extremely bad in any other case. As a result the standard metric in the statistics literature is the following.

Definition 8. For any estimator $\hat{\theta}_t$ we define its risk with t samples as,

$$R^t(\hat{\theta}) = \max_{p \in [0, 1]} E_p[|\hat{\theta}_t - p|]$$

Now our goal is clear, we need to design estimators $\hat{\theta}_t$ such that the risk $R^t(\hat{\theta}_t)$ is as small as possible. For example one can easily prove that the risk rate of the *sample mean estimator* ($\hat{\theta}_t = \frac{\sum_{i=1}^t X_i}{t}$) $R^t(\hat{\theta}_t) \leq \frac{1}{\sqrt{2t}}$. Obviously the next question is whether we can find another (probably more complex estimator) that achieves risk smaller than that of $\frac{1}{\sqrt{2t}}$.

Theorem 7. Any estimator $\hat{\theta}$ has risk $R^t(\hat{\theta}) \geq \frac{1}{16\sqrt{t}}$

The above theorem can be easily derived by standard techniques (e.g. Le Cam or Fano method) that have been developed in the statistics literature to provide lower bounds on the risk of estimators. Especially, proving this lower bound for the Bernoulli case requires a simple application of these methods. The above theorem also tells us that the risk of the *sample mean estimator* is up to constant factors optimal. Are we done?

Theorem 8. There exists an estimator $\hat{\theta}$ such that for all $p \in [0, 1]$,

$$E_p[|\hat{\theta}_t - p|] = O\left(\frac{1}{\sqrt{t}}\right)$$

Can this rate be improved? For example, is there an estimator $\hat{\theta}$ such that for all $p \in [0, 1]$, $E_p[|\hat{\theta}_t - p|] = O\left(\frac{1}{t}\right)$?

It seems that the latter question can be answered negatively using Theorem 13. Unfortunately, this is not the case. The difference is subtle but very important. The definition of the risk $R^t = \max_{p \in [0, 1]} E_p[|\hat{\theta}_t - p|]$, allows the maximum to be attained in a p that depends on t (e.g. $p = 1/2 + 1/t$). As a result, Theorem 13 does not provide us with information about the above question.

Theorem 9. For any estimator $\hat{\theta}$ there exists a fixed $p \in [0, 1]$ such that

$$E_p[|\hat{\theta}_t - p|] = \Omega\left(\frac{1}{t}\right)$$

To the best of our knowledge this question is not answered in the statistics literature. However, the same question has been asked for more the estimator of more general and complex distribution. These results ensure the above statement for estimators that satisfy certain properties (unbiased estimators, locally regular, e.t.c.). In this section we provide a theorem that holds for any estimator and indicates that estimators that don't satisfy the above statement are very unlikely to exist.

Theorem 10. For any Bernoulli estimator $\hat{\theta}$, for all $[a, b] \subseteq [0, 1]$,

$$\lim_{t \rightarrow \infty} t \int_a^b E_p[|\hat{\theta}_t - p|] dp = +\infty$$

In case that statement .. is not true than there exists an estimator $\hat{\theta}$ that simultaneously satisfies the following two properties:

1. for all $p \in [0, 1]$, $\lim_{t \rightarrow \infty} t E_p[|\hat{\theta}_t - p|] = 0$
2. for all $[a, b] \subseteq [0, 1]$, $\lim_{t \rightarrow \infty} t \int_a^b E_p[|\hat{\theta}_t - p|] dp = +\infty$

We conjecture that there does not exist a function that satisfies both of these properties. Even in case that this is not true, such a function has a very degenerate form making it very difficult to be the risk function of an estimator.

6.2 Final Lower Bound

Theorem 11. Let a graph oblivious protocol P such that for any instance I , $E[||x^t - x^*||_\infty] \leq f(I)g(t)$, where $f : I \mapsto \mathbb{R}_+$ and $g : \mathbb{N} \mapsto \mathbb{R}_+$. Then $\lim_{t \rightarrow \infty} t f(t) \neq 0$.

Proof. We use the protocol P to construct an estimator $\hat{\theta}$ and then we use the theorem ... For any $k, n \in \mathbb{N}$ with $k \leq n$ we construct the following instance.

- A star graph with $n + 1$ nodes
- k leaves have $a_i = 1$ and $s_i = 1$
- $n - k$ leaves have $a_i = 1$ and $s_i = 0$
- the central node has $a_{n+1} = 1/2$ and $s_{n+1} = 0$
- directed edges from the central node to all the leaves.

Observe that for all $i \neq n + 1$, $x_i^* = s_i$ and $x_{n+1}^* = \frac{k}{2n}$. For all $i \neq n + 1$, $x_i^t = P(1, 1, t)$ if $s_i = 1$ and $x_i^t = P(0, 1, t)$ if $s_i = 0$. Notice that these are deterministic functions with respect to t and are independent of the values of k, n . Now, $x_{n+1}^t = P(P(y_1, 1, 1), P(y_2, 1, 2), \dots, P(y_t, 1, t), 1/2, 0, t)$, where $x_i \in [0, 1]$. As a

result, $x_{n+1}^t = g(y_1, \dots, y_t)$, where $g : \{0, 1\}^t \mapsto [0, 1]$. As a result, the estimator $\hat{\theta}$ with $\hat{\theta}_t = \frac{x_t^t}{2}$ is valid estimator. Now let $p = k/n$ we have that

$$\begin{aligned} E_p[||x^t - x^*||] &\geq E_p[|x_{n+1}^t - \frac{p}{2}|] \\ &= E_p[|\frac{\hat{\theta}_t}{2} - \frac{p}{2}|] \\ &= \frac{1}{2} E_p[|\hat{\theta}_t - p|] \end{aligned}$$

Finally, for all $p \in [0, 1]$ $E_p[|\hat{\theta}_t - p|] \leq 2f(I_{k,n})g(t) = f'(p)g(t)$. By theorem .., $\lim_{t \rightarrow \infty} t f(t) \neq 0$. ■

7 Lower bound: Graph Oblivious Protocols

7.1 Our Results for Graph Oblivious

In the previous section, we have seen that in the unlabelled communication model, a simple update rule acheives an error rate $E[||x^t - x^*||_\infty] \leq \sqrt{\log n} \frac{(\log t)^2}{\sqrt{t}}$ when $a \geq 1/2$. This rate is good enough from the perspective of the size of the graph, but very inefficient from the perspective of the steps needed to acheive a certain accuracy. More precisely, in order to acheive accuracy $\epsilon > 0$, $O\left(\frac{\log n}{\epsilon^2}\right)$ steps are needed. On the other hand, in the labelled case another simple update rule needs $O\left(d^2 \log \frac{1}{\epsilon}\right)$. The terms $\log n$ and d^2 are comparable, however $\log \frac{1}{\epsilon}$ and $\frac{1}{\epsilon^2}$ are of different order of magnitude. In this section, we investigate whether this gap between these two cases can be closed or is impact of a generic defeciency of the unlabelled communication model.

Definition 9. A graph oblivious update rule F , is a sequence of function $\{f_t\}_{t=1}^\infty$ where $f_t : [0, 1]^{t+2} \mapsto [0, 1]$.

These restricted class of update rules is very natural since it is very reasonable to assume that the update rule is not aware of the graph structure of network and must work for any structure. Formally, this is imposed in Definition 9 since $F = \{f_t\}_{t=1}^\infty$ is selected prior to knowledge of G . In the next definition, we define the error of an update rule F for an instance I .

Definition 10. Let the graph oblivious update rule $F = \{f_t\}_{t=1}^\infty$. For an instance $I = (G, s, a)$ we define the following process. At round t :

- Each agent i has a state $x_i(t) \in [0, 1]$
- For each agent i , the oracle O_G randomly picks an agent $j \in N_i$ and returns the value $x_j(t-1)$ to i . Let O_i^t denotes this value.
- Each agent i updates her value as follows: $x_i(t) = f_t(O_i^1, \dots, O_i^t, s_i, a_i)$

The error of the update rule F , for the instance I is denoted by $E_I[||x(t) - x^*||_\infty]$, where x^* is the equilibrium point of I .

For example in the update rule ??, $f_t(O_1, \dots, O_t, s_i, a_i) = (1 - a_i) \frac{\sum_{j=1}^t O_j}{t} + a_i s_i$ and we have shown that for any instance I with $\alpha > 1/2$, $E_I[||x(t) - x^*||_\infty] \in O\left(\frac{(\log t)^2}{\sqrt{t}}\right)$. We investigate whether the exists a graph oblivious update rule that for gatantees a better asymptotic rate for all instances I . We conjecture that the following statement holds.

Theorem 12. For any graph oblivious update rule, $F = \{f_t\}_{t=1}^\infty$ there exists an instance I such that

$$E_I[\|x_t - x^*\|_\infty] \in \Omega\left(\frac{1}{t}\right)$$

Although we were not able to prove the above claim, in the next section we provide a theorem that indicates that the existence of such a graph oblivious update rule is highly unlikely. In section .., we prove that such an update rule must satisfy simultaneous two technical properties that excludes any reasonable update rule. We believe that our difficulty is of technical nature and we leave the proof of the above statement as an open problem.

7.2 Our results for No regret

We have an agent that acts at the following simple setting. At every time t the agent chooses $x_t \in [0, 1]$. Then, he receives a function $f_t : [0, 1] \rightarrow [0, 1]$. The cost incurred by the agent for round t is $f_t(x_t)$. The goal of the player is to pick x_t in a way that minimizes the total cost for all rounds. We emphasize the fact that the agent has to choose x_t before he sees f_t , otherwise the problem becomes trivial. Formally:

Definition 11. Let K be the set of all functions $f : [0, 1] \rightarrow \mathbb{R}$. Then, a protocol is a sequence of functions $F_t : K^{t-1} \rightarrow \mathbb{R}$, where $x_t = F_t(f_1, \dots, f_{t-1})$

In order to measure the performance of the agent, we define the following quantity:

Definition 12. If an agent runs the protocol for T rounds, we define the regret of the protocol as:

$$\sup_{\{f_1, \dots, f_T\} \subseteq K} \left\{ \sum_{t=1}^T f_t(x^t) - \min_{x \in [0, 1]} \left\{ \sum_{t=1}^T f_t(x) \right\} \right\}$$

This quantity should ideally be a sublinear function of T , since this means that on average the cost incurred by the algorithm is close to the best decision in hindsight. A protocol P is said to be *no regret* if the regret is $o(T)$. In the next definition, we define the error of an update rule F for an instance I .

Definition 13. Let the no regret update rule $F = \{F_t\}_{t=1}^\infty$. For an instance $I = (G, s, a)$ we define the following process. At round t :

- Each agent i has a state $x_i(t) \in [0, 1]$
- Each agent i updates her value as follows: $x_i(t) = F_t(f_1, \dots, f_{t-1})$
- For each agent i , the oracle O_G randomly picks an agent $j \in N_i$ and let $O_i^t = j$. Then, the oracle returns the following function:

$$f_t(x) = \alpha (x - s_i)^2 + (1 - \alpha) (x - x_j(t-1))^2$$

The error of the update rule F , for the instance I is denoted by $E_I[\|x(t) - x^*\|_\infty]$, where x^* is the equilibrium point of I .

For example, if we have the protocol:

$$F_t(f_1, \dots, f_{t-1}) = \operatorname{argmin}_{x \in [0, 1]} \sum_{i=1}^t f_i(x)$$

then it is easy to see that

$$x_i(t) = (1 - a_i) \frac{\sum_{j=1}^{t-1} x_{O_i^j}(t)}{t} + a_i s_i$$

which is just the weighted sample mean.

7.3 Estimating Bernoulli Distributions

The problem definition is the following:

- There exists an unknown parameter $p \in [0, 1]$
- We receive i.i.d. observations X_i , where $X_i \sim B(p)$
- We output an estimate value $\hat{p} = \hat{\theta}_t(X_1, \dots, X_t)$ according to some function $\hat{\theta}_t : \{0, 1\}^t \mapsto [0, 1]$.

The goal is to appropriately select $\hat{\theta}_t$ such that \hat{p} and p are relatively close. For example a very known selection is

$$\hat{\theta}_t(X_1, \dots, X_t) = \frac{\sum_{i=1}^t X_i}{t}$$

which is the sample mean estimator

Definition 14. An estimator $\hat{\theta}$ is a collection of functions $\{\hat{\theta}_t\}_{t=1}^\infty$, where $\hat{\theta}_t : \{0, 1\}^t \mapsto [0, 1]$.

As the number of samples grows we would expect that an efficient estimator $\hat{\theta}$ produces more and more accurate values for p . The following definition provides us the metric to evaluate the quality of an estimator $\hat{\theta}$.

Definition 15. For any estimator $\hat{\theta}_t$ we define

$$E_p[|\hat{\theta}_t(X_1, \dots, X_t) - p|] = \sum_{(x_1, \dots, x^t) \in \{0, 1\}^t} |\hat{\theta}_t(x^1, \dots, x^t) - p| \cdot p^{\sum_{i=1}^t x^i} (1-p)^{t - \sum_{i=1}^t x^i}$$

The quantity $E_p[|\hat{\theta}_t(X_1, \dots, X_t) - p|]$ is the expected distance of the estimated value $\hat{\theta}_t$ from the parameter p , when the distribution of the samples is $B(p)$. To simplify notation we also denote it as $E_p[|\hat{\theta}_t - p|]$.

Notice that $E_p[|\hat{\theta}_t - p|]$ does not serve as a good metric since an efficient estimator $\hat{\theta}$ must guarantee good estimates for any parameter p . Consider an estimator $\hat{\theta}$ that outputs $1/2$ for all t . In this case, $E_p[|\hat{\theta}_t - p|] = |p - 1/2|$ meaning that this estimator performs optimally when $p = 1/2$, but extremely bad in any other case. As a result the standard metric in the statistics literature is the following.

Definition 16. For any estimator $\hat{\theta}_t$ we define its risk with t samples as,

$$R^t(\hat{\theta}) = \max_{p \in [0, 1]} E_p[|\hat{\theta}_t - p|]$$

Now our goal is clear, we need to design estimators $\hat{\theta}_t$ such that the risk $R^t(\hat{\theta}_t)$ is as small as possible. For example one can easily prove that the risk rate of the *sample mean estimator* ($\hat{\theta}_t = \frac{\sum_{i=1}^t X_i}{t}$) $R^t(\hat{\theta}_t) \leq \frac{1}{\sqrt{2t}}$. Obviously the next question is whether we can find another (probably more complex estimator) that achieves risk smaller than that of $\frac{1}{\sqrt{2t}}$.

Theorem 13. Any estimator $\hat{\theta}$ has risk $R^t(\hat{\theta}) \geq \frac{1}{16\sqrt{t}}$

The above theorem can be easily derived by standard techniques (e.g. Le Cam or Fano method) that have been developed in the statistics literature to provide lower bounds on the risk of estimators. Especially, proving this lower bound for the Bernoulli case requires a simple application of these methods. The above theorem also tells us that the risk of the *sample mean estimator* is up to constant factors optimal. Are we done?

Theorem 14. There exists an estimator $\hat{\theta}$ such that for all $p \in [0, 1]$,

$$E_p[|\hat{\theta}_t - p|] = O\left(\frac{1}{\sqrt{t}}\right)$$

Can this rate be improved? For example, is there an estimator $\hat{\theta}$ such that for all $p \in [0, 1]$, $E_p[|\hat{\theta}_t - p|] = O\left(\frac{1}{t}\right)$?

It seems that the latter question can be answered negatively using Theorem 13. Unfortunately, this is not the case. The difference is subtle but very important. The definition of the risk $R^t = \max_{p \in [0, 1]} E_p[|\hat{\theta}_t - p|]$, allows the maximum to be attained in a p that depends on t (e.g. $p = 1/2 + 1/t$). As a result, Theorem 13 does not provide us with information about the above question.

Theorem 15. *For any estimator $\hat{\theta}$ there exists a fixed $p \in [0, 1]$ such that*

$$E_p[|\hat{\theta}_t - p|] = \Omega\left(\frac{1}{t}\right)$$

To the best of our knowledge this question is not answered in the statistics literature. However, the same question has been asked for more the estimator of more general and complex distribution. These results ensure the above statement for estimators that satisfy certain properties (unbiased estimators, locally regular, e.t.c.). In this section we provide a theorem that holds for any estimator and indicates that estimators that don't satisfy the above statement are very unlikely to exist.

Theorem 16. *For any Bernoulli estimator $\hat{\theta}$, for all $[a, b] \subseteq [0, 1]$,*

$$\lim_{t \rightarrow \infty} t \int_a^b E_p[|\hat{\theta}_t - p|] dp = +\infty$$

In case that statement .. is not true than there exists an estimator $\hat{\theta}$ that simultaneously satisfies the following two properties:

1. for all $p \in [0, 1]$, $\lim_{t \rightarrow \infty} t E_p[|\hat{\theta}_t - p|] = 0$
2. for all $[a, b] \subseteq [0, 1]$, $\lim_{t \rightarrow \infty} t \int_a^b E_p[|\hat{\theta}_t - p|] dp = +\infty$

We conjecture that there does not exist a function that satisfies both of these properties. Even in case that this is not true, such a function has a very degenerate form making it very difficult to be the risk function of an estimator.

7.4 Reductions from update rules to Bernoulli Estimators

Lemma 11. *For any no regret protocol P , there exists a Bernoulli estimator $\hat{\theta}$ such that for all $q \in Q([0, 1])$, there exists an instance I_q such that*

$$E_q[|\hat{\theta}_t - q|] \leq 2E_{I_q}[||x^t - x^*||_\infty]$$

Proof. At round t , the graph oblivious protocol P has selected the functions $F_t(f_1, \dots, f_{t-1})$. An estimator $\hat{\theta}$ receives the samples Y_1, \dots, Y_t and outputs the value $\hat{\theta}_t(Y_1, \dots, Y_t) \in [0, 1]$. By Lemma (12) the opinions of neighbours are functions of Y_1, \dots, Y_{t-2} . We define:

$$g_t(Y_1, \dots, Y_t) = g_{Y_t}^t(Y_1, \dots, Y_t)$$

We consider the estimator $\hat{\theta}_t = \frac{1}{2} F_t(g_1(Y_1), \dots, g_t(Y_1, \dots, Y_t))$ and observe that $\hat{\theta}_t$ is a function $\{0, 1\}^t \mapsto [0, 1]$.

Now for any $q \in Q([0, 1])$, construct an appropriate instance I_q such that $E_q[|\hat{\theta}_t - q|] \leq 2E_{I_q}[||x^t - x^*||_\infty]$.

Since $q = \frac{k}{n}$ some $k, n \in \mathbb{N}$, the following instance I_q with $n + 1$ agents:

- A central agent with $s_c = 0$ and $a_c = 1/2$.
- Directed edges from the central agent to all the other agents.

- k agents with $s_i = 0$ and $a_i = 1$
- $n - k$ agents with $s_i = 1$ and $a_i = 1$

Notice that in this instance for all $i \neq c$, $x_i^* = s_i$ and $x_c^* = \frac{q}{2}$. We show that $E_q[|\hat{\theta}_t - q|] \leq 2E_{I_q}[||x^t - x^*||_\infty]$. At round t , if the oracle returns to the center agent the value $g_t(1, 1)$ of a 1-agent, then $Y_t = 1$ otherwise $Y_t = 0$. As a result, $\mathbf{P}[Y_t = 1] = q$ and

$$\begin{aligned} E_{I_q}[||x^t - x^*||_\infty] &\geq E_{I_q}[|x_c^t - x_c^*|] \\ &= E_q[|\frac{\hat{\theta}_t}{2} - \frac{q}{2}|] \end{aligned}$$

■

Lemma 12. Suppose the no regret protocol $F = \{F_i\}_{i=1}^t$ runs in the topology defined in the previous lemma. We define the sequence b_1, b_2, \dots, b_t such that $b_i = s$, where s is the s_i of the neighbour of center that was selected at time i . Then, for every time i there are functions g_0^i, g_1^i such that for all neighbours j with $s_i = 0$: $x_j(i) = g_0^i(b_1, \dots, b_{i-2})$ and for all neighbours j with $s_i = 1$: $x_j(i) = g_1^i(b_1, \dots, b_{i-2})$.

Proof. We are going to prove the statement inductively for i . If $i = 1$, then each agent hasn't seen any function yet and should make a choice for $x(1)$. By the definition of F_1 , all the agents select the same value, so $g_0^1 = g_1^1 = F_1$. We verified the base case.

Suppose the statement holds for all times $i \leq k$. In particular, it holds for $i = k$, so there exist g_0^k, g_1^k that are functions of b_1, \dots, b_{k-2} that satisfy the proposition. Since all neighbouring values up to time k are functions of b_1, \dots, b_{k-2} and since the functions that the center node receives are determined by the opinions of the neighbours and the choice at round $k - 1$, we get that $x_c(k)$ is a function of b_1, \dots, b_{k-1} . At the end of round k , each agent learns one more function based on the opinion of one of its neighbours. Observe that every neighbour of the center has only the center as neighbour. As a result, all nodes with $s_i = 0$ receive the same function ϕ_1 and all nodes with $s_i = 1$ receive the same function ϕ_2 . We notice that ϕ_1 and ϕ_2 are uniquely determined by $x_c(k)$ and thus of b_1, \dots, b_{k-1} .

Now, at time $k + 1$ first each agent evaluates the new opinion, based on the new function he received. Since the function that each node received is determined by his s_i and $x_c(k)$, it follows that all nodes that have the same s_i have the same value, which is a function of b_1, \dots, b_{k-1} . The inductive step is now complete. ■

8 A Graph Aware Protocol

Let P be a discrete distribution over $[d]$, and let S_1, \dots, S_t be t i.i.d samples drawn from P , i.e. $(S_1, \dots, S_t) \sim P^t$. The empirical distribution \hat{P}_t is the following estimator of the density of P .

$$\hat{P}_N(A) = \frac{\sum_{i=1}^N \mathbb{1}_{[S_i \in A]}}{N}, \quad (16)$$

where $A \subseteq [d]$. In words, \hat{P}_N simply counts how many times the value i appeared in the samples S_1, \dots, S_N . We will use the following version of the classical Vapnik and Chervonenkis inequality.

Lemma 13. Let \mathcal{A} be a collection of subsets of $\{0, \dots, d\}$ and Let $S_{\mathcal{A}}(N)$ be the Vapnik-Chervonenkis shatter coefficient, defined by

$$S_{\mathcal{A}}(N) = \max_{x_1, \dots, x_t \in [d]} |\{\{x_1, \dots, x_t\} \cap A : A \in \mathcal{A}\}|.$$

Then

$$\mathbf{E}_{P^N} \left[\max_{A \in \mathcal{A}} |\hat{P}_t(A) - P(A)| \right] \leq 2\sqrt{\frac{\log 2S_{\mathcal{A}}(N)}{N}}$$

An interesting question is whether $\text{poly}(1/\varepsilon)$ rounds are necessary when protocols have access to the unlabeled oracle. We show that this is not true by showing a protocol that needs $\ln^2(1/\varepsilon)$ rounds to be within error ε from the solution x^* . Our protocol is graph-aware in the sense that it is allowed to depend on the graph G . Therefore, we show that without any constraints on the protocols it is hard to prove strong lower bounds for our problem. The problem in this case is that all agents could agree to stop updating their public opinions for enough rounds so that everybody learns, with high probability, exactly the average of the opinions of their neighbors. Having the exact average they can perform a step of the original update of the FJ-model which results in a similar and fast convergence rate. Of course if the opinions of the neighbors are all different we can still use an argument similar to that of the protocol of section COUPONS COLLECTOR. Since some neighbors may share opinions we have to come up with a different solution for this problem. Let $L \leq d$ be the number of different opinions and k_i be the number of neighbors who have opinion ξ_i . Then the average of the opinions of the neighbors is $(\sum_{i=1}^L k_i \xi_i) / d_i$. Instead of using the average like we did in section SAMPLE MEAN, we will instead find *exactly* the frequencies k_i and then compute their average exactly with high probability. We next describe our protocol. Since it is the same for all agents, For simplicity we describe it for agent i .

Algorithm 2 Graph Aware Update Rule

```

1:  $x_i(0) \leftarrow s_i, t \leftarrow 0$ .
2: for  $l = 1, \dots, \infty$  do
3:    $\varepsilon \leftarrow 1/2^l$ .
4:   for  $t = 1, \dots, \ln(1/\varepsilon)$  do
5:     Keep a map  $h$  from  $[0, 1]$  to  $[d_i]$  and array of counters  $A$  of length  $d_i$  and an array  $B$  of length  $M_1$ .
6:      $M_1 = O(\ln(1/\varepsilon)), M_2 = O(d^2 \ln(d))$ 
7:     for  $j = 1, \dots, M_1$  do
8:       for  $k = 1, \dots, M_2$  do
9:         Call the unlabeled oracle  $O_i^u(t)$  to get an opinion  $X_i$ .
10:        if  $X_i$  is not in  $h$  then
11:          Insert  $X_i$  to  $h$ .
12:        else
13:           $A(h(X_i)) \leftarrow A(h(X_i)) + 1$ .
14:         $t \leftarrow t + 1$ 
15:      Divide all entries of  $A$  by  $M_2$ .
16:      Round all entries of  $h$  to the closest multiple of  $1/d$ .
17:       $B(j) = \sum_{i=1}^{d_i} A(i)$ .
18:     $x_i(t) \leftarrow \text{maj}_j B(j)$ .
```

Theorem 17. *The update rule 2 for $a > 1/2$ achieves convergence rate*

$$\mathbb{E} [\|x_t - x^*\|_\infty] \leq C e^{-\sqrt{t}/(d\sqrt{\log d})},$$

where C is a universal constant.

Proof. According to the update rule 2 all agents fix their opinions $x_i(t)$ for $M_1 \times M_2$ rounds. To estimate the sum of the opinions each agent estimates the frequencies k_j/d_i . Since the neighbors have at most d_i different opinions we can think of the opinions as natural numbers in $[d_i]$. The oracle returns the opinion of a random neighbor and therefore the samples X_i returned by the oracle O_i^u are drawn from a discrete distribution P supported on $[d_i]$. The probability $P(j)$ of the j -th opinion is the number of neighbors having this opinion k_j/d_i . To learn the probabilities $P(j)$ using samples from P . Letting $\mathcal{A} = \{\{1\}, \{2\}, \dots, \{d_i\}\}$ we have from Lemma 13 that

$$\mathbb{E}_{P^m} \left[\max_{j \in [d_i]} |\hat{P}_m(j) - P(j)| \right] \leq 2\sqrt{\frac{\log 2d_i}{m}},$$

since $S_{\mathcal{A}} \leq d_i$. Therefore, an agent can draw $m = 100d^2 \log(2d)$ to learn the frequencies k_j/d_i within expected error $1/(5d)$. Notice now that the array A after line 15 corresponds to the empirical distribution of equation (16). Notice that if the agents have estimations of the frequencies k_j/d_i with error smaller than $1/d$ then by rounding them to the closest multiple of $1/d_i$ they learn the frequencies exactly. By Markov's inequality we have that with probability at least $4/5$ the rounded frequencies are exactly correct. By standard Chernoff bounds we have that if the agents repeat the above procedure $\ln(1/\delta)$ times and keep the most frequent of the answers $B(j)$ then they will obtain the correct answer with probability at least $1 - \delta$. From Lemma ?? we have that to achieve error ε we need $\log(1/\varepsilon)/(1/(1 - \alpha))$ rounds. Since we need all nodes to succeed at computing the exact averages for $\log(1/\varepsilon)/\log(1/(1 - \alpha))$ rounds we have from the union bound that for $\delta < \frac{\varepsilon \ln(1/(1 - \alpha))}{n \ln(1/\varepsilon)}$, with probability at least $1 - \varepsilon$ the error is at most ε . For the expected error after $T = O(d^2 \log d (\log(n/\varepsilon) + \log \log(1/\varepsilon) + \log \log(1 - \alpha)))$ rounds we have that $\mathbb{E} [\|x_T - x^*\|_\infty] = O(\varepsilon)$. ■

References

- [BKO11] David Bindel, Jon M. Kleinberg, and Sigal Oren. How bad is forming your own opinion? In *IEEE 52nd Annual Symposium on Foundations of Computer Science, FOCS 2011, Palm Springs, CA, USA, October 22-25, 2011*, pages 57–66, 2011.
- [GS14] Javad Ghaderi and R. Srikant. Opinion dynamics in social networks with stubborn agents: Equilibrium and convergence rate. *Automatica*, 50(12):3209–3215, 2014.
- [ZLZ17] Yichi Zhou, Jialian Li, and Jun Zhu. Identify the nash equilibrium in static games with random payoffs. In *Proceedings of the 34th International Conference on Machine Learning, ICML 2017, Sydney, NSW, Australia, 6-11 August 2017*, pages 4160–4169, 2017.